

Research Design: A Multi-Modal XAI Deep Learning framework for pneumonia diagnosis: Evaluating diagnostic performance and clinical trust over diverse populations.

By

[Stiaan Megit](#)

0211205065082

Literature Review is submitted in partial fulfilment of the requirements for the Bachelor
of

Computing Degree

at Belgium Campus iTversity

Supervisor: [Ely Ondo-Ekogha](#)

Date of submission: 30/08/2026

1. Introduction:

Artificial Intelligence (AI), particularly deep learning, has advanced Computer-Aided Diagnosis (CAD) for pneumonia detection from chest X-ray (CXR) images. Developments in convolutional neural networks, transfer learning, multi-modal learning, Explainable AI (XAI), and Human-in-the-Loop (HITL) approaches have improved diagnostic capability while supporting greater transparency and clinician involvement.

However, the literature review identified persistent limitations in diagnostic accuracy, interpretability, generalisability, and clinician-centred decision support. Existing approaches frequently address these challenges independently rather than within an integrated diagnostic framework. This study therefore proposes a **multi-modal XAI deep learning framework for pneumonia diagnosis** that combines transfer learning, patient metadata, XAI, and HITL decision support.

The primary contribution is the design and evaluation of this integrated artefact using a **Design Science Research (DSR)** methodology. The DSR approach provides a structured process for identifying the research problem, defining solution objectives, developing the artefact, demonstrating its operation, and evaluating its effectiveness while supporting methodological rigour and research contribution.

2. Justification of the Research Design:

2.1 Nature of the Research Problem:

Although deep learning, multi-modal learning, XAI, and HITL approaches have advanced AI-assisted pneumonia diagnosis, existing research continues to treat diagnostic performance, interpretability, generalisability, and human involvement largely as separate challenges. The research problem therefore extends beyond evaluating an existing model and requires the design of an integrated solution addressing these challenges collectively.

2.2 Why Design Science Research is Appropriate:

DSR is appropriate because it focuses on designing, developing, demonstrating, and evaluating artefacts that address real-world problems while generating research knowledge. Unlike approaches focused primarily on synthesising existing knowledge or testing predefined hypotheses, DSR is solution oriented. The primary artefact developed in this study is a multi-modal XAI pneumonia diagnostic framework integrating DenseNet-121 transfer learning, metadata fusion, Grad-CAM explanations, and HITL review.

The **Peppers DSR process model** provides a structured sequence for identifying the problem, defining objectives, designing and developing the artefact, demonstrating its application, and evaluating its outcomes. This aligns directly with the objectives of the study and supports systematic development and evaluation.

2.3 Alignment of Design Science Research with this Study:

This approach aligns with the principles of DSR methodology, in which the methodology's iterative nature supports the continuous refinement and evaluation of the proposed novel framework throughout its development process. Its rigorous process guides the design, development, and experimental evaluation of the artefacts against established performance criteria, enabling the systematic improvements of the proposed diagnostic framework.

3. Design Science Research Process:

This study follows the Peffers DSR methodology to develop and evaluate a lightweight multi-modal XAI framework for pneumonia diagnosis using CXR images and patient metadata. The artefact integrates DenseNet-121 transfer learning, metadata fusion, Grad-CAM, HITL review, subgroup evaluation, and preprocessing comparison through six experiments:

- **E1:** Image-only diagnostic baseline.
- **E2:** Multi-modal diagnostic model using CXR images and patient metadata.
- **E3:** Explainability using Grad-CAM.
- **E4:** Human-in-the-Loop review.
- **E5:** Generalisability across demographic and imaging subgroups.
- **E6:** Effect of preprocessing choices on diagnostic performance.

3.1 Problem Identification and Motivation:

Pneumonia diagnosis from CXR images remains challenging because of overlapping radiographic manifestations and variation across patients and imaging conditions. Although deep learning can provide effective automated classification, image-only approaches may omit relevant patient information, learn dataset-specific visual patterns, and provide limited insight into their predictions. These limitations motivate an integrated approach combining clinical metadata, visual explanations, human review, subgroup evaluation, and preprocessing analysis.

3.2 Objectives of the Solution:

The objective is to design and evaluate a lightweight multi-modal diagnostic framework that addresses predictive performance, interpretability, human involvement, and subgroup behaviour. DenseNet-121 transfer learning establishes the image-only baseline before age, sex, and view position are incorporated as metadata. Grad-CAM provides visual explanations, while HITL review enables human agreement, disagreement, and correction. Subgroup analysis examines age, sex, and view position, while CLAHE is evaluated as an alternative preprocessing strategy. Together, these components extend the artefact beyond standalone classification toward integrated diagnostic decision support.

3.3 Design and Development of the Artefact:

The developed artefact is a DSR research prototype designed to investigate a multi-modal approach to binary pneumonia classification using the NIH ChestX-ray14 dataset. The framework is structured as a sequential pipeline consisting of **data preparation, image classification, metadata fusion, explainability, human review, and evaluation components**. The design progressively extends an image-only baseline into a multi-modal diagnostic decision-support framework and subsequently evaluates its behaviour under different conditions.

- **Data Preparation:** Prepares the NIH ChestX-ray14 dataset for binary pneumonia classification by retaining cases labelled Pneumonia and No Finding. Patient metadata consisting of age, sex, and radiographic view position is encoded alongside the corresponding CXR images. To prevent images from the same patient from occurring across different data partitions, the dataset is divided at patient level into 70% training, 15% validation, and 15% test sets, using a fixed random seed. This produces a held-out test set for subsequent evaluation while reducing the risk of patient-level data leakage.
- **E1 - Image-only Baseline:** Establishes the baseline diagnostic architecture using an ImageNet-pretrained DenseNet-121 model with CXR images as the sole input. The architecture uses the DenseNet-121 backbone to extract hierarchical image features, followed by a task-specific projection and classification head for binary pneumonia prediction. This experiment establishes the reference architecture against which the contribution of additional patient metadata can subsequently be assessed.
- **E2 - Multi-modal Diagnostic Model:** Extends the E1 image-only architecture by incorporating structured patient metadata consisting of age, sex, and radiographic view position. The CXR image is processed through the DenseNet-121 image encoder, while the metadata is processed through a separate metadata encoder. The resulting image and metadata representations are then fused before being passed to the final classification layer, allowing information from both visual and structured clinical sources to contribute to the pneumonia prediction.
- **E3 - Explainable Artificial Intelligence:** Extends the multi-modal model with an explainability component using Gradient-weighted Class Activation Mapping (Grad-CAM). Grad-CAM is applied to the final convolutional layer of the DenseNet-121 backbone to generate spatial heatmaps highlighting image regions associated with the model's prediction. The component is designed to support examination of model behaviour across both correct and incorrect predictions by providing a visual representation of the image features contributing to the model output.
- **E4 - Human-in-the-Loop:** Incorporates human oversight into the diagnostic workflow by allowing a reviewer to assess the model's predictions together with their corresponding Grad-CAM explanations. The reviewer can record agreement or disagreement with the prediction, provide a corrected label where necessary, and optionally provide a supporting comment. This component positions the model as a


decision-support system in which human reviewers retain the ability to challenge and correct automated predictions.

- **E5 - Generalisability and Subgroup Evaluation:** Evaluates whether the multi-modal model maintains consistent behaviour across different patient and imaging subgroups within the held-out NIH ChestX-ray14 test set. The analysis considers age, sex, and radiographic view, with the test population divided into predefined demographic and imaging groups. Diagnostic performance is subsequently assessed within each subgroup to identify differences in model behaviour. This represents within-dataset subgroup generalisability rather than external validation, as all evaluations are performed using the NIH ChestX-ray14 test set.
- **E6 - Preprocessing Comparison:** Extends the evaluation of the multi-modal model by investigating whether the image preprocessing strategy influences diagnostic performance. Two configurations are compared: the baseline preprocessing pipeline and a CLAHE-based preprocessing pipeline. The same multi-modal architecture, dataset partitions, training configuration, and experimental conditions are maintained between the two configurations, allowing the preprocessing strategy to serve as the primary experimental variable.

The E1–E4 components form the core progressive diagnostic artefact, progressing from image-only classification to multi-modal prediction, visual explanation, and human oversight. E5 and E6 extend the artefact through subgroup analysis and preprocessing comparison. Collectively, these experiments evaluate the framework across diagnostic performance, multi-modal information integration, interpretability, human oversight, subgroup behaviour, and preprocessing sensitivity.

3.4 Demonstration:

The artefact was demonstrated using an end-to-end experimental workflow built on the NIH ChestX-ray14 dataset which operationalises the multi-modal diagnostic framework.



It progresses from image and patient metadata preparation to model inference, representative case selection, explainability, and human review. The held-out test set was treated as unseen data, ensuring that the demonstration reflected the application of the developed artefact to previously unused cases.

Following patient-level partitioning, CXR images were resized and normalised, while age, sex, and view position were numerically encoded and synchronised with the images. E1 implemented the image-only DenseNet-121 baseline, while E2 incorporated the encoded metadata. Both models were trained on the training set and evaluated on the validation set, after which the selected E2 checkpoint was used for held-out test-set inference.


For test-set inference, the trained E2 model was evaluated in inference mode without image augmentation. Each test image was processed together with its corresponding metadata to generate a sigmoid-based pneumonia probability, which was converted into a binary prediction using the configured classification threshold. The resulting predictions, probabilities, and ground-truth labels were used to select four representative cases: a true positive (TP), true negative (TN), false positive (FP), and false negative (FN). These cases were subsequently processed through E3, where Grad-CAM heatmaps and overlays were generated to visualise the image regions contributing to each prediction. Quantitative measures, including mean activation and central activation ratio, were also recorded to support the visual assessment of the explanations.

Finally, E4 demonstrated the human-in-the-loop component by presenting the selected cases, model probabilities, predictions, and Grad-CAM overlays to a reviewer through the implemented review interface. The reviewer recorded agreement or disagreement, provided a corrected label where required, and could add a supporting comment. E5 and E6 then extended the artefact beyond the core diagnostic workflow. E5 evaluated the multi-modal model across predefined age, sex, and radiographic-view subgroups to identify differences in diagnostic behaviour across patient and imaging characteristics. E6 compared the baseline preprocessing pipeline with a CLAHE-based alternative to determine whether the choice of image preprocessing influenced model performance under otherwise consistent experimental conditions.

3.5 Evaluation:


Evaluation was performed using the predefined patient-level dataset partitions, with the held-out test set reserved for final assessment. The held-out test set was used for final evaluation, with a probability threshold of 0.50 applied during training and validation and 0.75 used for test-set classification following validation threshold analysis.

E1 established the image-only DenseNet-121 baseline, while E2 incorporated age, sex, and radiographic view position. As shown in **Diagram**, E2 improved accuracy, precision, specificity, F1-score, and ROC-AUC, although recall and PR-AUC decreased slightly. This provided a direct comparison of image-only and multi-modal diagnostic performance.



E3 assessed explainability using representative TP, TN, FP, and FN cases. Grad-CAM heatmaps and activation measures demonstrated differing spatial patterns across predictions, as illustrated in **Figure 3.X**. This evaluated the artefact's explanation capability rather than the clinical validity of the explanations.

Figure 3.X: Grad-CAM visualisations for representative true-positive, true-negative, false-positive, and false-negative cases



E4 assessed human oversight using the four representative cases. The reviewer agreed with two predictions and disagreed with two; both model errors were corrected, resulting in 100% reviewer accuracy for the selected cases. The evaluation therefore demonstrated the functionality of the human-in-the-loop component.

Table 3.X: E4 human-in-the-loop evaluation

Evaluation measure	Result
Cases reviewed	4
Model accuracy	50%
Reviewer agreement with model	50%
Reviewer accuracy	100%
Model errors	2
Errors corrected	2 (100%)

E5 evaluated the E2 model across age, sex, and radiographic-view subgroups. As shown in **Table 3.X**, performance varied most noticeably between PA and AP views, while reduced sensitivity was observed in the

60+ age group. This demonstrated differences in model behaviour across subgroups within the held-out test set.

Table 3.X: E5 subgroup evaluation of the multi-modal model

Subgroup	n	Accuracy	Sensitivity	Specificity	F1	ROC-AUC
All test	9,326	0.9239	0.3026	0.9394	0.1627	0.7815
Age <40	3,205	0.9332	0.3086	0.9494	0.1894	0.7968
Age 40–59	4,071	0.9167	0.3438	0.9306	0.1630	0.7722
Age 60+	2,050	0.9234	0.2157	0.9415	0.1229	0.7819
Female	4,072	0.9197	0.2600	0.9363	0.1372	0.8034
Male	5,254	0.9271	0.3359	0.9419	0.1834	0.7642
PA	5,969	0.9648	0.1683	0.9785	0.1393	0.7123
AP	3,357	0.8511	0.4094	0.8684	0.1722	0.7884

Finally, E6 compared baseline and CLAHE preprocessing under matched conditions. **Table 3.X** shows that baseline preprocessing achieved the higher validation ROC-AUC (0.7982 versus 0.7846), indicating that CLAHE did not improve performance in this experiment. Collectively, the evaluations examined the artefact's predictive, explanatory, interactive, subgroup, and preprocessing characteristics.

Table 3.X: E6 preprocessing comparison

Metric	Baseline	CLAHE	Difference
Epoch 1 ROC-AUC	0.7300	0.7285	−0.0015
Epoch 2 ROC-AUC	0.7975	0.7680	−0.0295
Epoch 3 ROC-AUC	0.7982	0.7846	−0.0136
Epoch 3 F1-score	0.1175	0.1102	−0.0073
Epoch 3 Sensitivity	0.7452	0.7019	−0.0433
Epoch 3 Specificity	0.7536	0.7516	−0.0020

3.6Communication:

4. Research Design Decisions:

5. Reliability, Validity and Rigor:

6. Ethical Considerations:

7. Limitations of the Research Design:

8. Conclusion:

3.6 Communication

The final DSR activity concerns communicating the artefact, its design decisions, evaluation results, and limitations to the research and practitioner communities.

The artefact is communicated through this research design paper, the documented experimental pipeline, quantitative evaluation metrics, Grad-CAM visualisations, HITL feedback records, subgroup analysis results, and preprocessing comparison results.

The implementation produces structured CSV outputs containing experimental metrics and case-level results, supporting an auditable record of the evaluation process.

The findings communicate several contributions. First, patient metadata can be integrated with a DenseNet-121 CXR classifier within a lightweight multi-modal architecture. Second, Grad-CAM can be incorporated to expose spatial characteristics associated with model predictions. Third, human review can provide a mechanism for detecting and correcting model errors. Fourth, subgroup evaluation demonstrates that aggregate diagnostic performance can conceal differences across demographic and imaging groups. Finally, preprocessing evaluation demonstrates that alternative image enhancement techniques should be empirically validated rather than assumed to improve performance.

The communication of these findings also identifies the boundaries of the artefact, particularly the lack of external validation, the limited HITL sample, the limited XAI case sample, and the absence of a large-scale clinical user study.

4. Research Design Decisions

Several methodological decisions were made to ensure that the research remained aligned with the DSR methodology and the objectives established in the previous papers.

Design Science Research as the Primary Methodology

DSR was selected because the primary contribution of the study is the **design and evaluation of an artefact**, rather than the observation of an existing phenomenon alone. The framework represents an engineered solution to the identified problem and can therefore be evaluated according to its functionality, performance, interpretability, and utility.

The six activities of Peffers et al. (2007) provide the structure through which the artefact was developed, demonstrated, evaluated, and communicated.

Experimental Evaluation

Experimental evaluation was selected as a supporting method within the DSR process. E1 and E2 provide a comparison between image-only and multi-modal classification. E3 and E4 extend the evaluation beyond predictive performance to interpretability and human involvement. E5 investigates subgroup behaviour, while E6 examines the effect of preprocessing choices.

This combination allows the artefact to be evaluated from multiple perspectives rather than through a single performance metric.

Transfer Learning

DenseNet-121 was selected as the underlying image architecture to utilise established deep visual representations without requiring the development of an entirely new convolutional architecture. Transfer learning also supports the objective of developing a lightweight research artefact within the available computational resources.

Multi-Modal Integration

Age, sex, and radiographic view were incorporated as metadata because these variables are available alongside the CXR image and provide contextual information that is not directly represented by image pixels.

The purpose was not to claim that these variables are inherently diagnostic, but to investigate whether contextual patient and acquisition information can contribute to model discrimination.

Grad-CAM

Grad-CAM was selected as the primary XAI technique because it provides spatially interpretable heatmaps that can be directly associated with CXR regions. It also has precedent in pneumonia research, including Arias-Londoño et al. (2020) and Sheu et al. (2023).

SHAP and LIME were not implemented in the final artefact. This was a deliberate scope decision rather than an indication that these techniques are unsuitable. Implementing multiple XAI techniques would have increased experimental complexity without necessarily improving the core DSR demonstration.

Human-in-the-Loop Design

HITL was incorporated because an automated diagnostic system should not necessarily operate as an autonomous replacement for clinical judgement. Instead, the framework demonstrates a decision-support approach in which humans can inspect, challenge, and correct model outputs.

Subgroup Evaluation

Subgroup evaluation was incorporated to investigate whether aggregate model performance was representative of different demographic and imaging groups.

Age, sex, and radiographic view were selected because these variables were already available within the NIH dataset and form part of the metadata used by the multi-modal model.

The subgroup analysis is therefore intended to identify performance variation rather than to claim comprehensive fairness validation.

Preprocessing Design

The preprocessing comparison was introduced to determine whether an alternative image enhancement strategy could improve diagnostic performance.

CLAHE was selected because it can enhance local image contrast and is commonly considered for medical imaging applications. However, the experiment was designed as an empirical comparison rather than assuming that contrast enhancement would improve classification.

The results ultimately demonstrated that the baseline preprocessing pipeline produced higher final ROC-AUC and PR-AUC than CLAHE under the tested conditions.

5. Reliability, Validity and Rigor

Several measures were applied to improve the reliability, validity, and methodological rigor of the research.

First, the E1 and E2 models used the same NIH ChestX-ray14 dataset structure and evaluation procedure, allowing the contribution of patient metadata to be examined under comparable conditions. The E2 model achieved a best validation ROC-AUC of 0.7982 and was subsequently evaluated on the held-out test set as part of E5.

Second, multiple diagnostic metrics were used rather than relying solely on accuracy. Accuracy, precision, recall, sensitivity, specificity, F1-score, ROC-AUC, and PR-AUC were recorded. This is particularly important because the evaluation dataset is highly imbalanced, with only 228 positive cases among 9,326 test images.

Consequently, the overall accuracy of 92.39% must be interpreted alongside the much lower precision of 11.13% and recall of 30.26%. The use of multiple metrics reduces the likelihood that conclusions will be driven by a single potentially misleading measure.

Third, patient-level splitting was used during dataset preparation. Patients were assigned to training, validation, and test partitions rather than randomly splitting individual images. This reduces the risk that images from the same patient appear across different partitions and artificially inflate performance.

Fourth, the XAI evaluation deliberately included TP, TN, FP, and FN cases. This reduces the risk of presenting only successful predictions and allows model explanation behaviour to be examined for both correct and erroneous outcomes.

Fifth, E5 expanded the evaluation beyond aggregate test performance by analysing age, sex, and radiographic-view subgroups. This provides additional evidence regarding the consistency of model behaviour within the evaluated population.

The subgroup results demonstrate why this analysis is important. For example, ROC-AUC ranged from 0.7123 for PA-view images to 0.8034 for female patients, while recall ranged from 16.83% for PA images to 40.94% for AP images.

Sixth, E6 provided a controlled preprocessing comparison. Both preprocessing configurations used the same multi-modal architecture and dataset structure, while the preprocessing strategy was changed. This enables the observed difference to be interpreted primarily in relation to the preprocessing choice.

The final E6 results showed that baseline preprocessing achieved a validation ROC-AUC of 0.7982 compared with 0.7846 for CLAHE. PR-AUC showed the same direction, with 0.0789 for baseline preprocessing compared with 0.0738 for CLAHE.

Seventh, the framework records model predictions, probabilities, metadata, explanation metrics, reviewer decisions, and corrected labels in structured output files. This creates an auditable experimental record and supports reproducibility.

Finally, the use of DSR provides methodological structure by explicitly connecting the identified problem, solution objectives, artefact, demonstration, evaluation, and communication activities.

Despite these measures, the study does not provide statistical evidence of clinical validity. The experiments should therefore be interpreted as evaluation of a research artefact rather than clinical validation of a deployable diagnostic system.

6. Ethical Considerations

The use of medical imaging and patient metadata introduces important ethical considerations. Although the study uses an existing research dataset rather than collecting new patient data, medical images and demographic attributes require careful handling.

The artefact was therefore designed as a research and decision-support system rather than an autonomous clinical diagnostic system. The predictions generated by the model should not be interpreted as medical diagnoses or as a replacement for qualified healthcare professionals.

Patient Privacy

The framework uses the de-identified NIH ChestX-ray14 dataset and limited metadata consisting of age, sex, and radiographic view. No directly identifying patient information is incorporated into the proposed diagnostic workflow.

Patient identifiers are used only as dataset-level identifiers necessary for patient-level data partitioning and evaluation.

Algorithmic Bias and Subgroup Differences

Algorithmic bias represents an important consideration because model performance can vary according to demographic and acquisition characteristics.

The E5 subgroup evaluation demonstrated measurable performance differences across age, sex, and radiographic view. For example, PA images achieved a ROC-AUC of 0.7123, whereas AP images achieved 0.7884. Female patients achieved a ROC-AUC of 0.8034 compared with 0.7642 for male patients.

These differences do not by themselves establish algorithmic discrimination or clinical unfairness, but they demonstrate why aggregate performance should not be treated as sufficient evidence of reliable behaviour across all populations.

The subgroup analysis therefore serves as an initial investigation of performance variation rather than a comprehensive fairness assessment.

Human Oversight

The incorporation of HITL provides an additional ethical safeguard because the framework positions the model as an aid to human decision-making rather than an unquestionable authority.

In the demonstrated cases, the reviewer was able to identify and correct both incorrect model predictions. However, the small sample and single-reviewer design prevent conclusions about clinical effectiveness.

Explainability

XAI introduces another ethical consideration. A Grad-CAM heatmap indicates image regions associated with a model prediction, but it does not establish that those regions represent a clinically verified pathological cause.

Consequently, Grad-CAM explanations should be used to support human inspection and model auditing rather than being treated as definitive clinical evidence.

Responsible Use

The framework should therefore be regarded as an experimental research artefact. Deployment in a clinical environment would require additional validation, appropriate regulatory assessment, security controls, clinical oversight, and evaluation by qualified medical professionals.

The results obtained in this study do not establish that the model is suitable for autonomous diagnosis or clinical deployment.

7. Limitations of the Research Design

Several limitations constrain the conclusions that can be drawn from this study.

First, the evaluation was conducted using the NIH ChestX-ray14 dataset. Although the E5 subgroup analysis evaluated differences across age, sex, and radiographic view, it remains a **within-dataset generalisability evaluation** rather than external validation. The study therefore does not establish whether the framework would perform similarly on images from other hospitals, imaging equipment, acquisition protocols, or patient populations.

Second, the positive class is substantially underrepresented. The E5 test set contained 228 positive cases compared with 9,098 negative cases. This imbalance contributes to the large difference between accuracy and positive-class metrics. Although the E2 model achieved 92.39% test accuracy, its precision was only 11.13% and recall was 30.26%.

Therefore, the high accuracy should not be interpreted as evidence of clinical-level pneumonia detection capability.

Third, the E5 subgroup analysis demonstrates that performance varies across different groups. In particular, PA-view images produced a ROC-AUC of 0.7123 and recall of only 16.83%, whereas AP-view images produced a ROC-AUC of 0.7884 and recall of 40.94%. These differences indicate that aggregate performance can conceal important variations in model behaviour.

However, the subgroup analysis is limited to the metadata variables available in the selected dataset. It does not constitute a comprehensive fairness or bias assessment across all clinically relevant demographic and socioeconomic characteristics.

Fourth, the XAI evaluation was intentionally limited to four representative cases. Although the inclusion of TP, TN, FP, and FN cases provides qualitative coverage of different prediction outcomes,

four cases are insufficient to establish the statistical reliability or generalisability of explanation behaviour.

Fifth, only Grad-CAM was implemented. Although SHAP, LIME, Grad-CAM++, and other explanation methods have been investigated in previous research, the current artefact does not provide a systematic comparison of different XAI techniques.

Sixth, the HITL evaluation involved only four reviewed cases and a single reviewer. Although both model errors in the selected cases were successfully corrected, the result cannot be generalised to clinical professionals or larger populations. The 100% reviewer correctness observed in the selected cases should therefore not be interpreted as evidence of clinical reliability or trust.

Seventh, E6 was limited to a comparison between the baseline preprocessing pipeline and CLAHE. Although the experiment found that baseline preprocessing produced a higher final validation ROC-AUC and PR-AUC than CLAHE, this does not establish that the baseline pipeline is universally optimal. Other preprocessing methods, image normalisation strategies, resolution choices, or augmentation techniques may produce different results.

Eighth, the E6 experiments were conducted using three training epochs because of the available computational resources. Although three epochs were sufficient to observe measurable differences between the preprocessing configurations, longer training could potentially alter the relative performance of the two approaches.

Ninth, the study does not constitute a clinical validation study. No prospective hospital trial, statistically meaningful clinician study, or comparison against radiologist performance was conducted.

Finally, the framework prioritised feasibility and demonstration within the available computational and time constraints. More extensive hyperparameter optimisation, external validation, larger-scale XAI evaluation, multiple XAI techniques, larger HITL studies, and additional demographic subgroup analysis could provide stronger evidence of effectiveness and generalisability.

8. Conclusion

This study designed and evaluated a lightweight multi-modal XAI deep learning framework for pneumonia diagnosis using Design Science Research. The resulting artefact integrates DenseNet-121-based image classification, patient metadata, Grad-CAM explanations, Human-in-the-Loop review, subgroup analysis, and preprocessing evaluation into a unified diagnostic decision-support workflow.

The E1 and E2 experiments established the progression from an image-only diagnostic model to a multi-modal architecture incorporating age, sex, and radiographic view. The E2 model achieved a best validation ROC-AUC of **0.7982**, demonstrating its ability to learn meaningful discriminatory patterns during model development.

Evaluation on the held-out NIH test set in E5 produced an overall accuracy of **92.39%**, specificity of **93.94%**, F1-score of **16.27%**, ROC-AUC of **0.7815**, and PR-AUC of **0.0858**. However, recall was only **30.26%** and precision was **11.13%**, demonstrating the limitations of relying on accuracy in a highly imbalanced diagnostic problem.

The E5 subgroup analysis further demonstrated that aggregate performance does not fully represent model behaviour. ROC-AUC ranged from **0.7123 for PA-view images** to **0.8034 for female patients**, while recall ranged from **16.83% for PA images** to **40.94% for AP images**. These findings demonstrate measurable variation across patient and imaging subgroups and support the inclusion of subgroup analysis when evaluating AI-based diagnostic systems.

The E3 XAI experiment demonstrated that Grad-CAM can expose differences in model activation between correct and incorrect predictions. The false-positive and false-negative cases demonstrated that prediction confidence and activation behaviour do not necessarily correspond directly to diagnostic correctness. This reinforces the value of examining model behaviour rather than relying exclusively on prediction probabilities.

The E4 HITL experiment further demonstrated that human review can provide a mechanism for identifying and correcting erroneous predictions. Both model errors in the selected four-case demonstration were successfully corrected by the reviewer. Nevertheless, the small sample size and single-reviewer design mean that these results should be interpreted only as a proof of concept rather than evidence of clinical trust.

Finally, E6 demonstrated that preprocessing choices can influence diagnostic performance. The baseline preprocessing strategy achieved a validation ROC-AUC of **0.7982** at the third epoch compared with **0.7846** for CLAHE. PR-AUC showed a similar pattern, with **0.0789** for baseline preprocessing compared with **0.0738** for CLAHE. Under the evaluated conditions, therefore, CLAHE did not provide an improvement over the baseline preprocessing pipeline.

Overall, the study demonstrates the feasibility of integrating **multi-modal learning, explainable AI, human oversight, subgroup evaluation, and preprocessing analysis within a single lightweight diagnostic framework**.

The primary contribution of the research is therefore not the achievement of state-of-the-art pneumonia classification accuracy. Instead, the contribution lies in the design and evaluation of an integrated DSR artefact that addresses several limitations identified in existing pneumonia CAD research.

The findings also demonstrate that evaluating an AI diagnostic framework requires more than a single aggregate accuracy measure. Diagnostic discrimination, positive-class detection, explanation behaviour, human correction, subgroup variation, and preprocessing sensitivity all provide important information about the behaviour and limitations of the artefact.

Future research should therefore focus on **external validation using independent datasets, larger clinician-based HITL studies, systematic comparison of XAI techniques including SHAP and LIME, broader subgroup and fairness analysis, optimisation of preprocessing and training strategies, and evaluation of explanation reliability**. Such work would provide stronger evidence regarding whether the proposed framework could ultimately support reliable clinical decision-making across diverse healthcare environments.

References

- Agarwal, S., 2022. COVID-19 & Pneumonia Detection using Multi Model Fusion. *Interantional Journal of engineering science*, 22 January.14(4).
- Bhatia, B., 2024. Evaluating Explainable Artificial Intelligence (XAI) techniques in chest radiology imaging through a human-centered Lens. 9 October.
- Confalonier, R., 2020. A historical perspective of explainable Artificial Intelligence. *WIREs Data Mining and Knowledge Discovery*, 19 October.11(1).
- Dominquez-Rodriquez, S., 2023. Testing the performance, adequacy, and applicability of an artificial intelligence model for pediatric pneumonia diagnosis. *Computer Methods and Programs in Biomedicine*, 10 December. Volume 242.
- Islam, M. N., 2023. Interpretable Differential Diagnosis of Non-COVID Viral Pneumonia, Lung Opacity and COVID-19 Using Tuned Transfer Learning and Explainable AI. *Healthcare*, 31 January.11(3).
- Kumar, R., 2025. Enhanced Multi-Model Deep Learning for Rapid and Precise Diagnosis of Pulmonary Diseases Using Chest X-Ray Imaging. *Diagnostics*, 22 January.15(3).
- Lo, S.-H., 2021. A novel interaction-based methodology towards explainable AI with better understanding of Pneumonia Chest X-ray Images. *Discover Artificial Intelligence*, 20 December.1(16).
- Mahamud, E., 2024. An explainable artificial intelligence model for multiple lung diseases classification from chest X-ray images using fine-tuned transfer learning. *Decision Analytics Journal*, 11 September. Volume 12.
- Mustapha, J. C., 2025. Advancing malware imagery classification with explainable deep learning: A state-of-the-art approach using SHAP, LIME and Grad-CAM. 28 may.
- Naz, Z., 2023. An Explainable AI-Enabled Framework for Interpreting Pulmonary Diseases from Chest Radiographs. *Cancers*, 3 January.15(1).
- Nia, N. G., 2023. Evaluation of artificial intelligence techniques in disease diagnosis and prediction. *Discover Artificial intelligence*, 30 January.3(5).
- Rauf, A., 2024. Exploring the current and prospective role of artificial intelligence in disease diagnosis. *Annals of Medicine & Surgery*, 4 January, Volume 86, pp. 943-949.
- Ren, H., 2021. Interpretable Pneumonia Detection by Combining Deep Learning and Explainable Models With Multisource Data. *IEEE Access*, 17 June. Volume 9.

Shah, M., 2022. Advancement of deep learning in pneumonia/Covid-19 classification and localization: A systematic review with qualitative and quantitative analysis. *MedNexus*, 26 September.

Sheu, R.-K., 2023. Interpretable Classification of Pneumonia Infection Using eXplainable AI (XAI-ICP). *IEEEEn Xplore*, 10 March. Volume 11.

Varkey, B., 2020. Principles of Clinical Ethics and Their Application to Practice. *Medical Principles and Practice*, 4 January.30(1).

Yazdani, E., 2025. Explainable artificial intelligence for pneumonia classification: Clinical insights into deformable prototypical part network in pediatric chest x-ray images. *Journal of Medical Imaging and Radiation Sciences*, 10 September.56(5).

Zhang, Y., 2022. Applications of Explainable Artificial Intelligence in Diagnosis and Surgery. *Diagnostics*, 2 December.12(2).